

▼ Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
import math
```

▼ Importing Dataset

```
dataset=pd.read_csv('car performance.csv')
dataset
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino
...
393	27.0	4	140.0	86	2790	15.6	82	1	ford mustang gl

▼ Finding missing data

```
dataset.isnull().any()
```

mpg	False
cylinders	False
displacement	False
horsepower	False

```
weight          False
acceleration    False
model year      False
origin          False
car name        False
dtype: bool
```

There are no null characters in the columns but there is a special character '?' in the 'horsepower' column. So we replaced '?' with nan and replaced nan values with mean of the column.

```
dataset['horsepower']=dataset['horsepower'].replace('?',np.nan)
```

```
dataset['horsepower'].isnull().sum()
```

```
0
```

```
dataset['horsepower']=dataset['horsepower'].astype('float64')
```

```
dataset['horsepower'].fillna((dataset['horsepower'].mean()),inplace=True)
```

```
dataset.isnull().any()
```

```
mpg          False
cylinders     False
displacement  False
horsepower    False
weight        False
acceleration  False
model year    False
origin        False
car name      False
dtype: bool
```

#Pandas dataframe.info() function is used to get a quick overview of the dataset.

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   float64
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model year       398 non-null   int64
7   origin           398 non-null   int64
8   car name         398 non-null   object
dtypes: float64(4), int64(4), object(1)
memory usage: 28.1+ KB
```

#Pandas describe() is used to view some basic statistical details of a data frame or a series of numbers.

```
dataset.describe()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year
count	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.165829	2970.424623	15.568090	76.010050
std	7.815984	1.701004	104.269838	38.298676	846.841774	2.757689	3.697627
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
50%	23.000000	4.000000	148.500000	92.000000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	125.000000	3608.000000	17.175000	79.000000

There is no use with car name attribute so drop it

```
#dropping the unwanted column.
dataset=dataset.drop('car name',axis=1)
```

```
#Pandas dataframe.corr() is used to find the pairwise correlation of all columns in the dataframe.
corr_table=dataset.corr()
corr_table
```


	mpg	cylinders	displacement	horsepower	weight	acceleration	model year
mpg	1.000000	-0.775396	-0.804203	-0.777501	-0.831741	0.420289	0.579267
cylinders	-0.775396	1.000000	0.950721	0.842437	0.896017	-0.505419	-0.348746
displacement	-0.804203	0.950721	1.000000	0.897082	0.932824	-0.543684	-0.370164
horsepower	-0.777501	0.842437	0.897082	1.000000	0.863990	-0.686436	-0.417081
weight	-0.831741	0.896017	0.932824	0.863990	1.000000	-0.417457	-0.306564
acceleration	0.420289	-0.505419	-0.543684	-0.686436	-0.417457	1.000000	0.288137
model year	0.579267	-0.348746	-0.370164	-0.417081	-0.306564	0.288137	1.000000

▼ Data Visualizations

Heatmap : which represents correlation between attributes

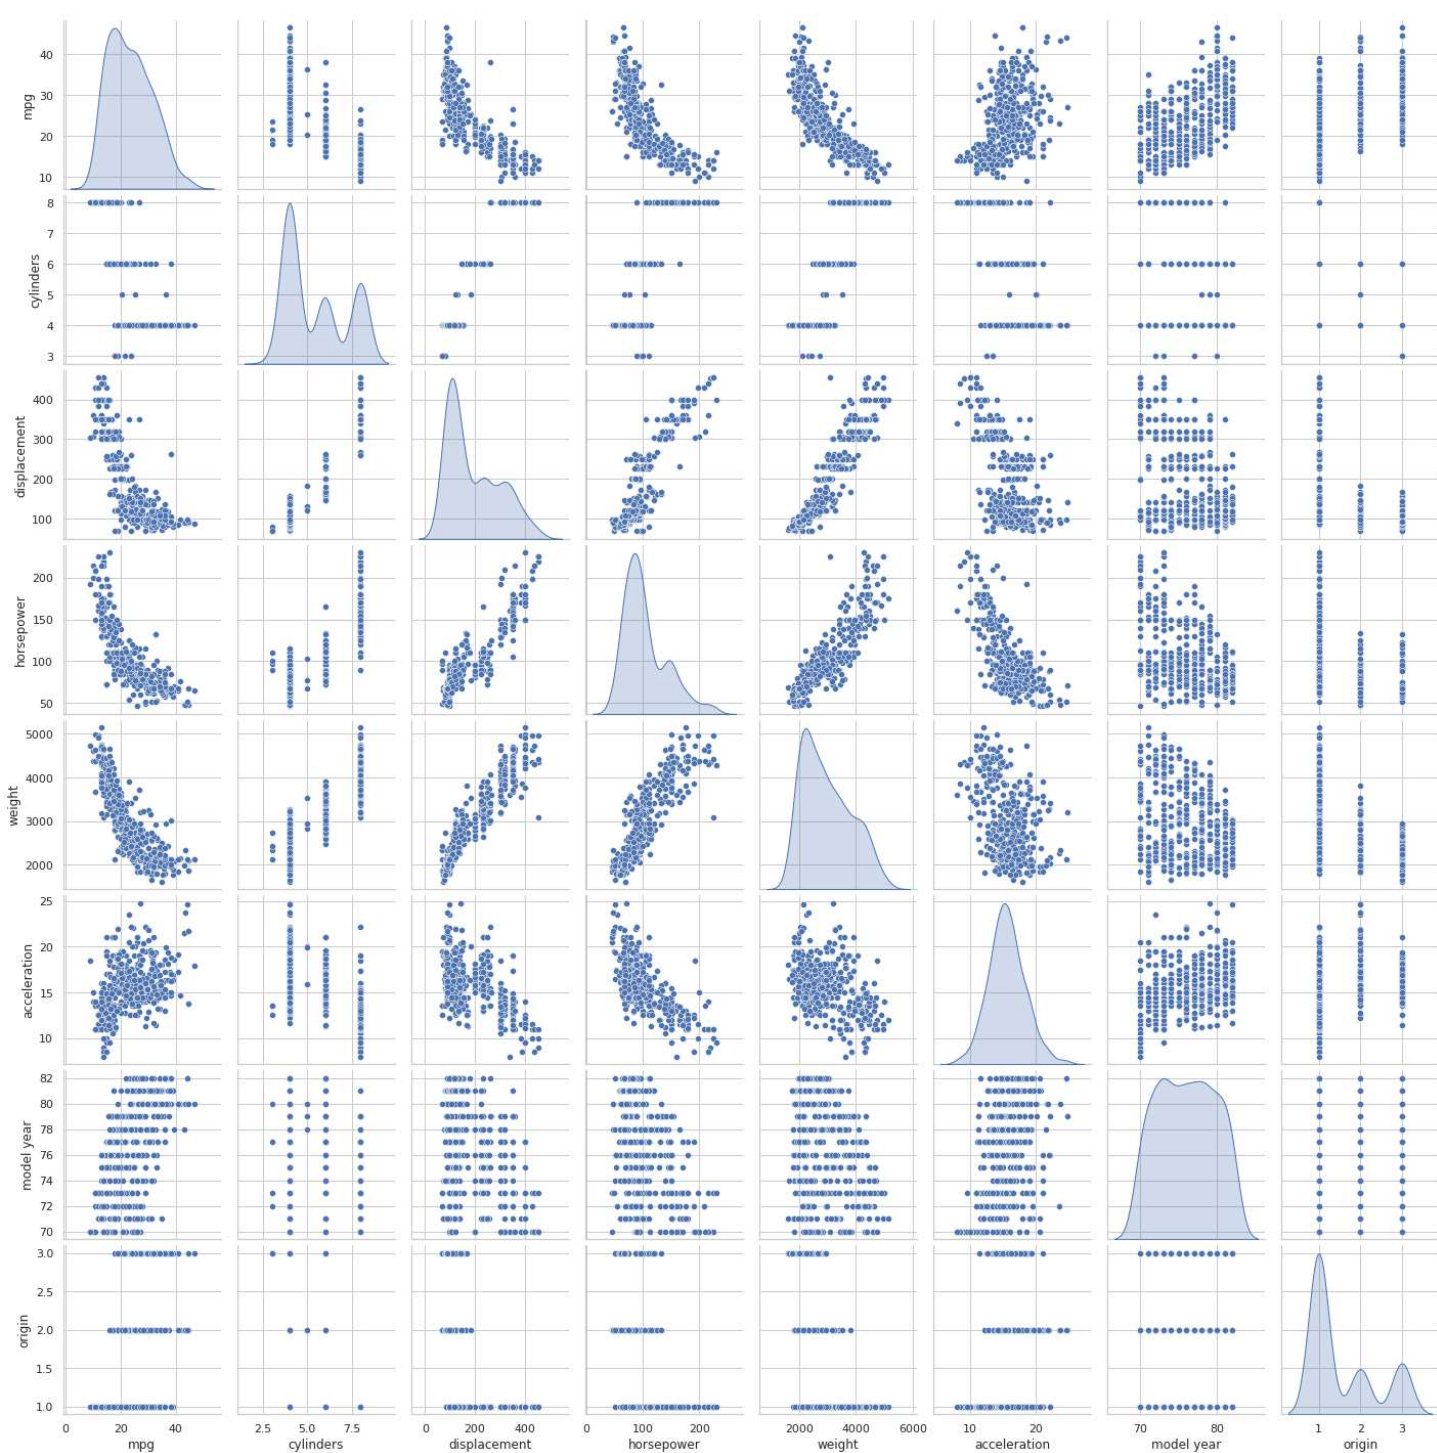
Double-click (or enter) to edit

```
#Heatmap is a way to show some sort of matrix plot,annot is used for correlation.
sns.heatmap(dataset.corr(),annot=True,linecolor='black', linewidths = 1)
fig=plt.gcf()
fig.set_size_inches(8,8)
```



Visualizations of each attributes w.r.t rest of all attributes


```
#pairplot represents pairwise relation across the entire dataframe.
sns.pairplot(dataset,diag_kind='kde')
plt.show()
```



Regression plots(`regplot()`) creates a regression line between 2 parameters and helps to visualize their linear relationships.

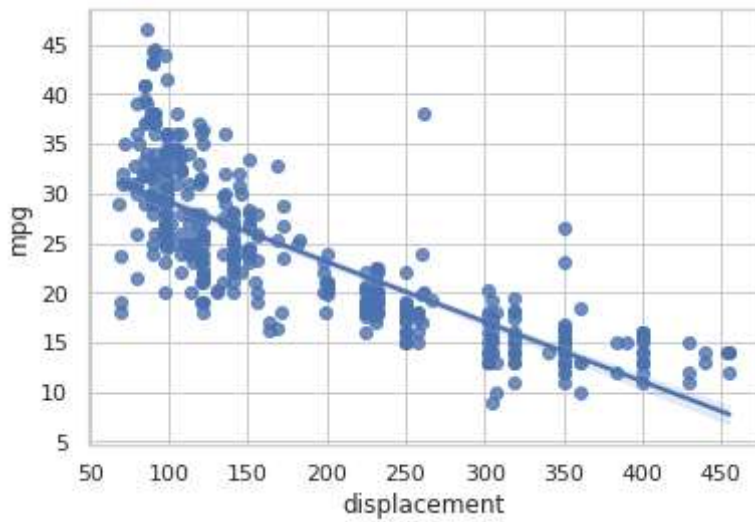
```
sns.regplot(x="cylinders", y="mpg", data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc037c3d0>
```




```
sns.regplot(x="displacement", y="mpg", data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc03b2490>
```



```
sns.regplot(x="horsepower", y="mpg", data=dataset)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bbfedd590>
```



```
sns.regplot(x="weight", y="mpg", data=dataset)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc0460190>  
sns.regplot(x="acceleration", y="mpg", data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc0b50dd0>
```




```
sns.regplot(x="model year", y="mpg", data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc04d38d0>
```



```
sns.regplot(x="origin", y="mpg", data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc06a8190>
```




```
sns.set(style="whitegrid")
```



```
sns.boxplot(x=dataset["mpg"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bc0ad5e10>
```



```
#finding the interquartilerange of mpg column and replacing the outliers with mean
```

```
q1=dataset['mpg'].quantile(0.25)
```

```
q3=dataset['mpg'].quantile(0.75)
```

```
iqr=q3-q1
```

```
lower_bound=q1-1.5*iqr
```

```
upper_bound=q3+1.5*iqr
```


```
dataset['mpg']=np.where(dataset['mpg']>upper_bound,dataset['mpg'].mean(),np.where(dataset['mpg']<lower_bound,dataset['mpg'].mean(),dataset['mpg']))
```

```
#boxplot for mpg column
```

```
sns.boxplot(dataset['mpg'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {\"x\": \"mpg\"}. This warning will disappear with Seaborn v0.12.0 and earlier.
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5bbd23b050>
```



Finding quartiles for mpg

The P-value is the probability value that the correlation between these two variables is statistically significant.

Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

- p-value is ≤ 0.001 : we say there is strong evidence that the correlation is significant.
- the p-value is ≤ 0.05 : there is moderate evidence that the correlation is significant.
- the p-value is ≤ 0.1 : there is weak evidence that the correlation is significant.
- the p-value is > 0.1 : there is no evidence that the correlation is significant

```
from scipy import stats
```

Cylinders vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Cylinders' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['cylinders'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

```
The Pearson Correlation Coefficient is -0.77764675219798 with a P-value of P = 7.858432477700
```



Conclusion:

Since the p-value is < 0.001 , the correlation between cylinders and mpg is statistically significant, and the coefficient of ~ -0.775 shows that the relationship is negative and moderately strong.

Displacement vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Displacement' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['displacement'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

```
The Pearson Correlation Coefficient is -0.805459196737545 with a P-value of P = 5.30475981248
```



Conclusion:

Since the p-value is < 0.1 , the correlation between displacement and mpg is statistically significant, and the linear negative relationship is quite strong (~ -0.809 , close to -1)

Horsepower vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['horsepower'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is -0.7785160459711739 with a P-value of P = 3.9815800890



Conclusion:


Since the p-value is < 0.001 , the correlation between horsepower and mpg is statistically significant, and the coefficient of ~ -0.771 shows that the relationship is negative and moderately strong.

Weght vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'weight' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['weight'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is -0.8334127823939774 with a P-value of P = 4.9210938207



Conclusion:

Since the p-value is < 0.001 , the correlation between weight and mpg is statistically significant, and the linear negative relationship is quite strong (~ -0.831 , close to -1)

Acceleration vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Acceleration' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['acceleration'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.4186346290331828 with a P-value of P = 2.55443395965



Conclusion:

Since the p-value is > 0.1 , the correlation between acceleration and mpg is statistically significant, but the linear relationship is weak (~ 0.420).

Model year vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Model year' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['model year'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.5776371269001912 with a P-value of P = 8.51521448588



Conclusion:

Since the p-value is < 0.001 , the correlation between model year and mpg is statistically significant, but the linear relationship is only moderate (~ 0.579).

Origin vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Origin' and 'mpg'.

```
pearson_coef, p_value = stats.pearsonr(dataset['origin'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.556374784654882 with a P-value of P = 1.006632653279

Conclusion:

Since the p-value is < 0.001 , the correlation between origin and mpg is statistically significant, but the linear relationship is only moderate (~ 0.563).

Ordinary Least Squares Statistics

```
test=smf.ols('mpg~cylinders+displacement+horsepower+weight+acceleration+origin',dataset).fit()
test.summary()
```

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.721
Model:	OLS	Adj. R-squared:	0.716
Method:	Least Squares	F-statistic:	168.0

Inference as in the above summary the p value of the acceleration is maximum(i.e 0.972) so we can remove the acc variable from the dataset

Df Residuals:	391	BIC:	2291.
----------------------	-----	-------------	-------

Seperating into Dependent and Independent variables

Intercept 43.8965 2.657 16.519 0.000 38.672 49.121

```
x=dataset[['cylinders','displacement','horsepower','weight','model year','origin']].values
x
```

```
array([[8.000e+00, 3.070e+02, 1.300e+02, 3.504e+03, 7.000e+01, 1.000e+00],
       [8.000e+00, 3.500e+02, 1.650e+02, 3.693e+03, 7.000e+01, 1.000e+00],
       [8.000e+00, 3.180e+02, 1.500e+02, 3.436e+03, 7.000e+01, 1.000e+00],
       ...,
       [4.000e+00, 1.350e+02, 8.400e+01, 2.295e+03, 8.200e+01, 1.000e+00],
       [4.000e+00, 1.200e+02, 7.900e+01, 2.625e+03, 8.200e+01, 1.000e+00],
       [4.000e+00, 1.190e+02, 8.200e+01, 2.720e+03, 8.200e+01, 1.000e+00]])
```

Kurtosis: 3.867 **Cond. No.** 4.00e+04

```
y=dataset.iloc[:,0:1].values
y
```

```
[22.8 ],
[23.5 ],
[30. ],
[39.1 ],
[39. ],
[35.1 ],
[32.3 ],
[37. ],
[37.7 ],
[34.1 ],
[34.7 ],
[34.4 ],
[29.9 ],
[33. ],

[34.5 ],
[33.7 ],
[32.4 ],
[32.9 ],
[31.6 ],
[28.1 ],
[30.7 ],
[25.4 ],
[24.2 ],
[22.4 ],
[26.6 ],
[20.2 ],
[17.6 ],
[28. ],
[27. ],
[34. ],
[31. ],
[29. ],
[27. ],
[24. ],
[22. ],
```

```
[25.    ],
[36.    ],
[37.    ],
[31.    ],
[38.    ],
[36.    ],
[36.    ],
[36.    ],
[34.    ],
[38.    ],
[32.    ],
[38.    ],
[25.    ],
[38.    ],
[26.    ],
[22.    ],
[32.    ],
[36.    ],
[27.    ],
[27.    ],
[44.    ],
[32.    ],
[28.    ],
[31.    ]])
```

▼ Splitting into train and test data.

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.1,random_state=0)
```

we are splitting as 90% train data and 10% test data

▼ random forest regressor

```
from sklearn.ensemble import RandomForestRegressor
```

```
rf= RandomForestRegressor(n_estimators=10,random_state=0)
model=rf.fit(x_train,y_train)
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DataConversionWarning: A column

```
y_pred=rf.predict(x_test)
y_pred
```

```
array([[14.55      , 25.89      , 14.7       , 21.4       , 18.       ,
        30.9       , 35.47145729, 22.71      , 15.8       , 26.19      ,
        34.34      , 33.52145729, 18.9       , 25.3       , 15.85      ,
        33.25      , 27.74      , 27.42      , 16.77      , 29.61291457,
        16.       , 25.44      , 23.84      , 20.76      , 32.52      ,
```

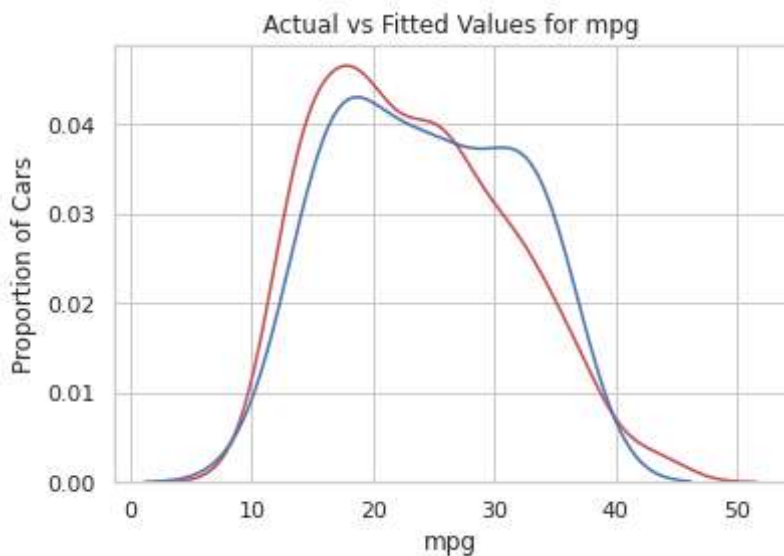
```
25.12      , 35.9      , 32.57      , 30.36      , 16.35      ,
19.87      , 31.       , 17.54      , 33.42      , 20.85      ,
23.8       , 19.17     , 16.15     , 35.7       , 11.5       ])
```

```
ax1 = sns.distplot(dataset['mpg'], hist=False, color="r", label="Actual Value")
sns.distplot(y_pred, hist=False, color="b", label="Fitted Values" , ax=ax1)
```

```
plt.title('Actual vs Fitted Values for mpg')
plt.xlabel('mpg')
plt.ylabel('Proportion of Cars')
```

```
plt.show()
plt.close()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot`
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot`
warnings.warn(msg, FutureWarning)
```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

Modal Evaluation

```
#importing necessary libraries to find evaluation of the model
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
import math
```

```
#mean squared error
MSE=mean_squared_error(y_test,y_pred)
print("MSE:",MSE)
```

```
MSE: 6.505788848703318
```

```
#Root mean squared error
```

```
RMSE=math.sqrt(MSE)
print("RMSE:",RMSE)
```

RMSE: 2.550644790774152


```
#checking the performance of the model using r2_score
r2=r2_score(y_test,y_pred)
print("R2_score:",r2)
```

R2_score: 0.9058760463516443

```
#Adjusted R square
Adjusted_R2=1-(1-r2*((x_test.shape[0]-1)/(x_test.shape[0]-x_test.shape[1]-1)))
print("Adjusted R2:",Adjusted_R2)
```

Adjusted R2: 1.0705807820519433

```
#plot for predicted and actual performance
plt.figure(figsize=(10,5))
plt.plot(y_pred[0:20])
plt.plot(np.array(y_test[0:20]))
plt.legend(["predicted","actual"])
plt.show()
```



```
accuracy = r2_score(y_test,y_pred)
accuracy
```

0.9058760463516443

```
import joblib
```

```
joblib.dump(model,'car performance')
```

['car performance']

```
siva=joblib.load('car performance')
```



```
siva.predict([[ '4', '97', '88', '2130', '71', '3' ]])
```

```
array([27.])
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 23:30

